
Interlude: Quick Review

Interlude: Quick Review

Review: Task Scheduling

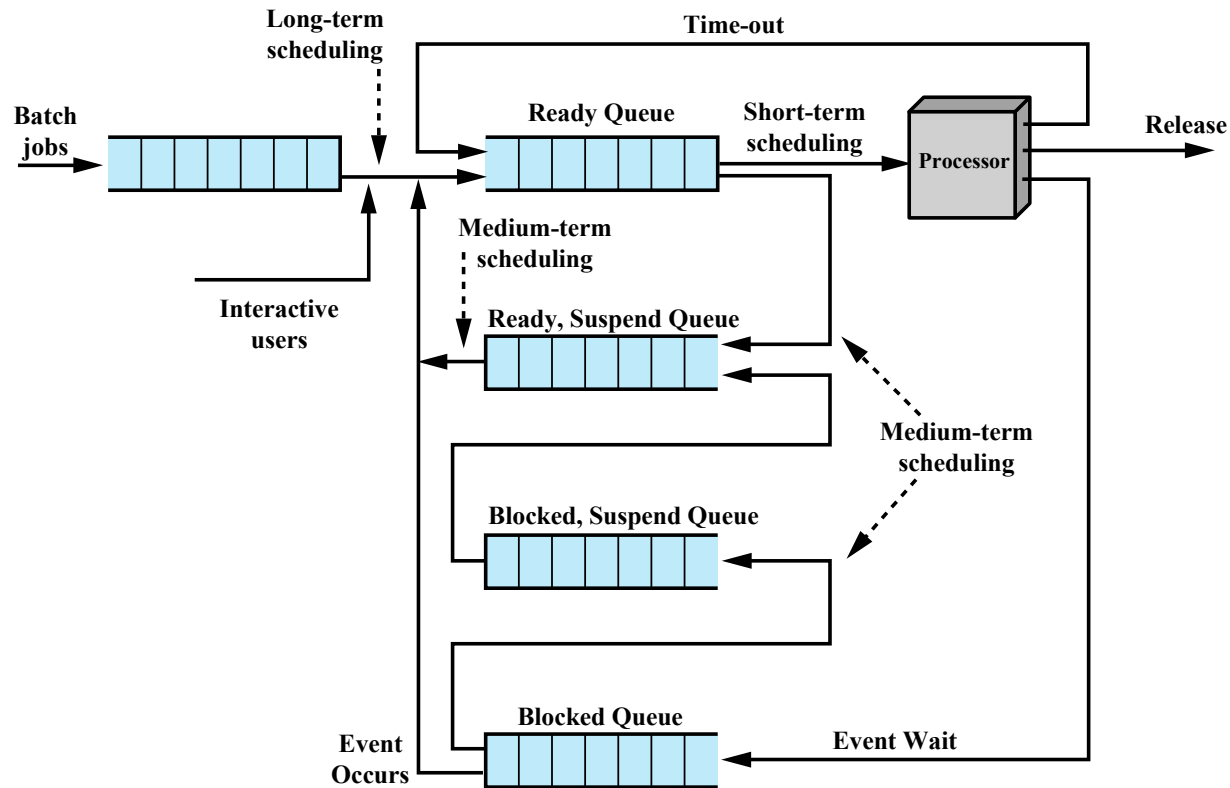


Figure 9.3 Queuing Diagram for Scheduling

Review: Bad Assumptions

All tasks:

1. ~~Run for the same amount of time~~ **FCFS**
2. ~~Arrive at the same time (roughly)~~ **SPN** shortest process next
3. ~~Once started, run to completion~~ **SRT** shortest remaining time
4. ~~Only use the CPU~~ **I/O**
5. ~~Have a known run-time.~~ **Exp Averaging**

Part 5: Feedback Queues

“Multi-Level Feedback Queue” (MLFQ) to be precise

Lower \rightarrow bigger #
higher \rightarrow smaller #

The Algorithm (Part 1)

1. don't estimate runtime

* smaller is more important

2. schedule preemptively, using multiple ready queues

$RQ_0 \rightarrow RQ_w$

a) Priority(task) = $N - RQ \#$ $0 \rightarrow$ highest priority

b) Priority(A) > Priority(B) \rightarrow schedule A

c) Priority(A) == Priority(B) \rightarrow Schedule using RR \rightarrow time slice

d) all tasks enter into RQ_0

e) On preemption, move the running task down a RQ

f)

g)

Visual





Figure 9.10 Feedback Scheduling

Simple Example

Task	Arrival Time	Service Time
A	$0 - \varepsilon$	4
B	$2 - \varepsilon$	6
C	$4 - \varepsilon$	4


$q=1$
3 RQs



...what happens to long(er) tasks?

Task	Arrival Time	Service Time
A	$0 - \epsilon$	4
B	$2 - \epsilon$	2
C	$4 - \epsilon$	2
D	$6 - \epsilon$	2
E	$8 - \epsilon$	2

a starves and gets
bad turnaround
time



The Algorithm (Part 2)


1. Don't estimate runtime
2. Schedule *preemptively* using *multiple* ready queues
 - a) $\text{Priority}(\text{task}) = n - \text{RQ \#}$
 - b) $\text{Priority}(A) > \text{Priority}(B) \rightarrow \text{schedule } A$
 - c) $\text{Priority}(A) = \text{Priority}(B) \rightarrow \text{schedule in RR}$
 - d) Tasks start in queue 0
 - e) On pre-emption, tasks move *down* a queue
 - f) increase the length of a time slice for lower queues
 - g)

each lower
q gets 2x
the ts
length
 $q = 2^i$

Redo



Task	Arrival Time	Service Time
A	$0 - \epsilon$	4
B	$2 - \epsilon$	2
C	$4 - \epsilon$	2
D	$6 - \epsilon$	2
E	$8 - \epsilon$	2

$$\begin{array}{r}
 B \quad A \quad 0 \\
 | \quad | \quad \hline
 B \quad A \quad 1 \\
 \hline
 2
 \end{array}$$



What about even *longer* tasks?

Task	Arrival Time	Service Time
A	$0 - \epsilon$	5
B	$2 - \epsilon$	2
C	$4 - \epsilon$	2
D	$6 - \epsilon$	2
E	$8 - \epsilon$	2



The Algorithm (Part 3)

1. Don't estimate runtime
2. Schedule *preemptively* using *multiple* ready queues
 - a) $\text{Priority}(\text{task}) = n - \text{RQ \#}$
 - b) $\text{Priority}(A) > \text{Priority}(B) \rightarrow$ schedule A
 - c) $\text{Priority}(A) = \text{Priority}(B) \rightarrow$ schedule in RR
 - d) Tasks start in queue 0
 - e) On pre-emption, tasks move *down* a queue
 - f) Longer time-slice for lower queues, e.g., $q = 2^i$
 - g) after a task has waited for S consecutive time units, promote to RQ_0

Analysis

- Does not require future knowledge
→ only based on the past
 - "May favor I/O bound tasks"
Fix: demote at the moment a proc has run for a total time slice
→ since only get demoted on pre-emption, not willfully giving up CPU
 - A good balance between response time and turnaround time
- ⇒ MLFQs form the basis for most UNIX schedulers

A decorative border consisting of three concentric, wavy lines. The outermost line is gray, the middle line is blue, and the innermost line is black. These lines form an irregular, cloud-like shape that frames the text.

Part 6: Wrapping Up

Adding Priority

Typically, four levels of priority:

1. System ← highest
2. Interactive
3. Normal
4. Batch

⇒ use priority queues

More Performance Metrics

Throughput: The # of tasks completed
per unit of time

↓
EXIT

Even more... (see book)

- Task deadlines
- Predictability
- Processor utilization
- Enforcing priorities
- Balancing resources
- Overhead

Formal Analysis

■ E : total execution time

how much time
has the task
gotter?

■ W : total waiting time

■ S : total service time

↑ either predict the future
or
Require future knowledge

Selection function:

which task is *selected* to be scheduled

Summary of Scheduling Algorithms

be able to fill this out for final → answer in ch 9

	FCFS	RR Round Robin	SPN	SRT	HRRN	MLFQ
Selection Function	$\max(w)$					
Preemptive?	X					
Starvation Possible?	X					
Response Time	poor					
Overhead	low					

med: 77
mean: 74

Midterm