

Chapters 11 & 12

Persistence

Part 0: Introduction

Persistent Storage

a device which stores information permanently, even through a power loss

→ Disk, SSDs, magnetic tape
↳ hard drive

→ Not RAM

Crux of the problem: How does the OS manage a persistent device?

Files and Directories

File: a linear array of bytes each of which you can r or w
→ a low level name (inode number)

Directory: a list of tuples: $\left(\begin{matrix} \text{user} \\ \text{readable} \\ \text{name} \end{matrix}, \begin{matrix} \text{inode} \\ \text{number} \end{matrix} \right)$

➤ also have an inode number

➤ so we can nest directories → directory tree

➤ The root directory: "/" in Unix

Example Directory Tree

need a separator:

"/" - sep in Unix

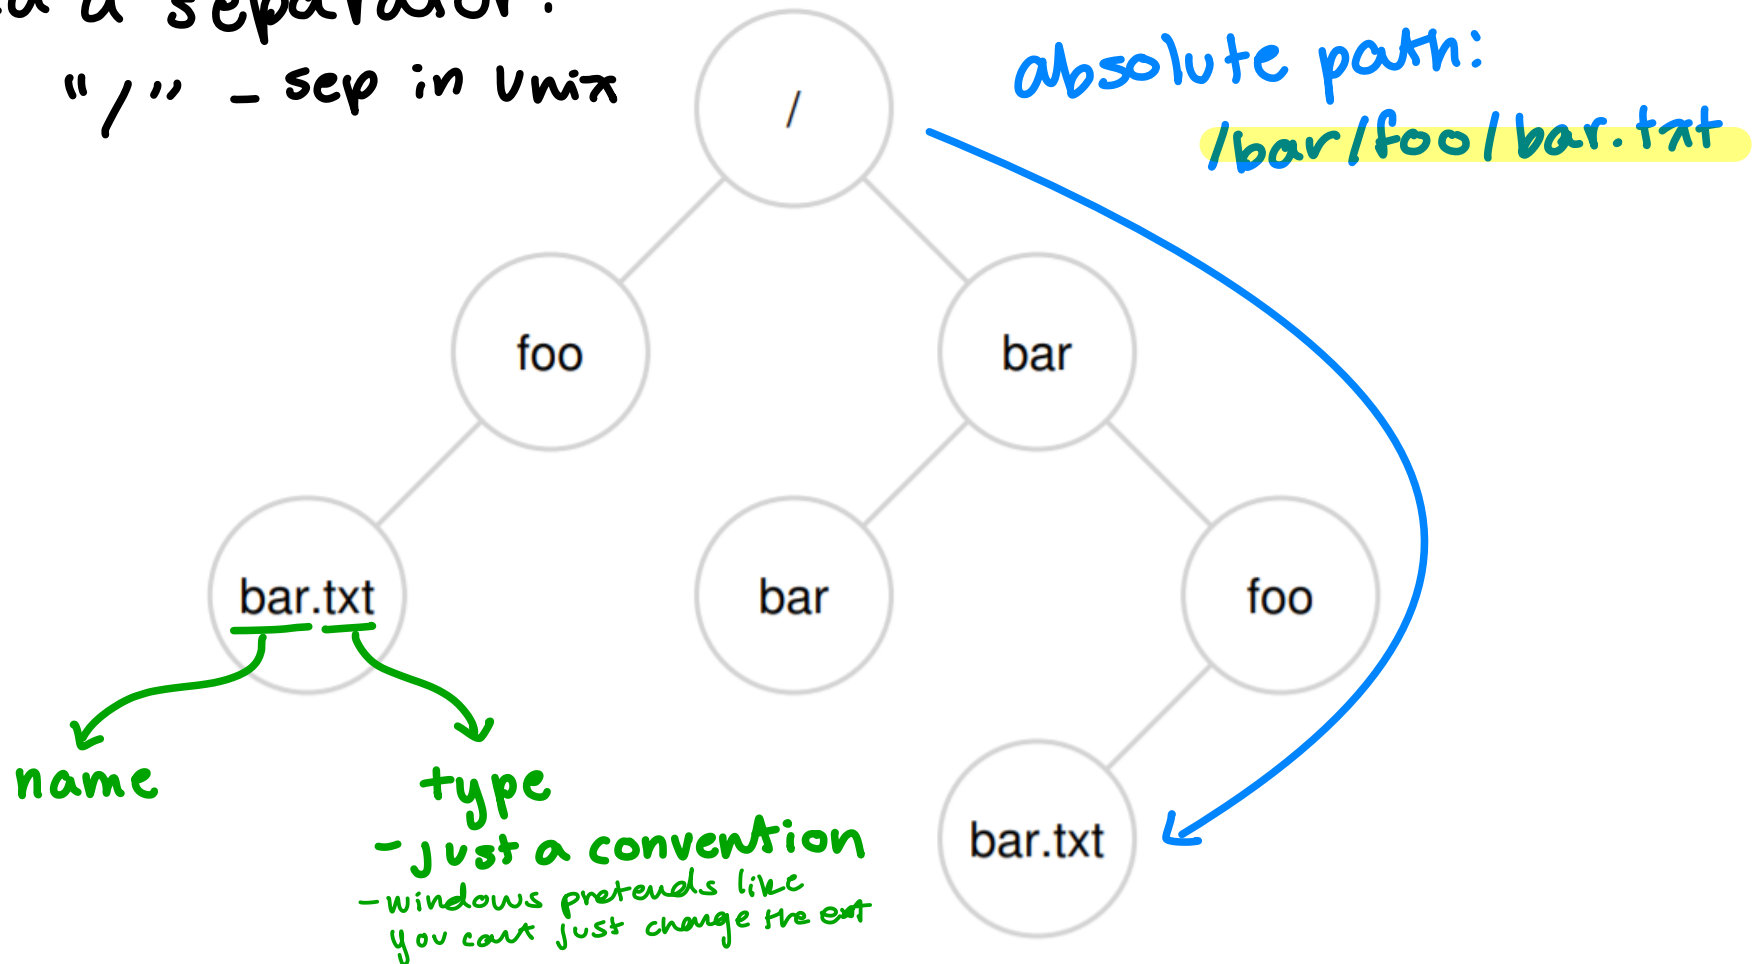


Figure 39.1: An Example Directory Tree

(Unix) File System Interface

➤ `open()`

➤ `read()`

➤ `write()`

➤ `close()`

➤ `dup()`

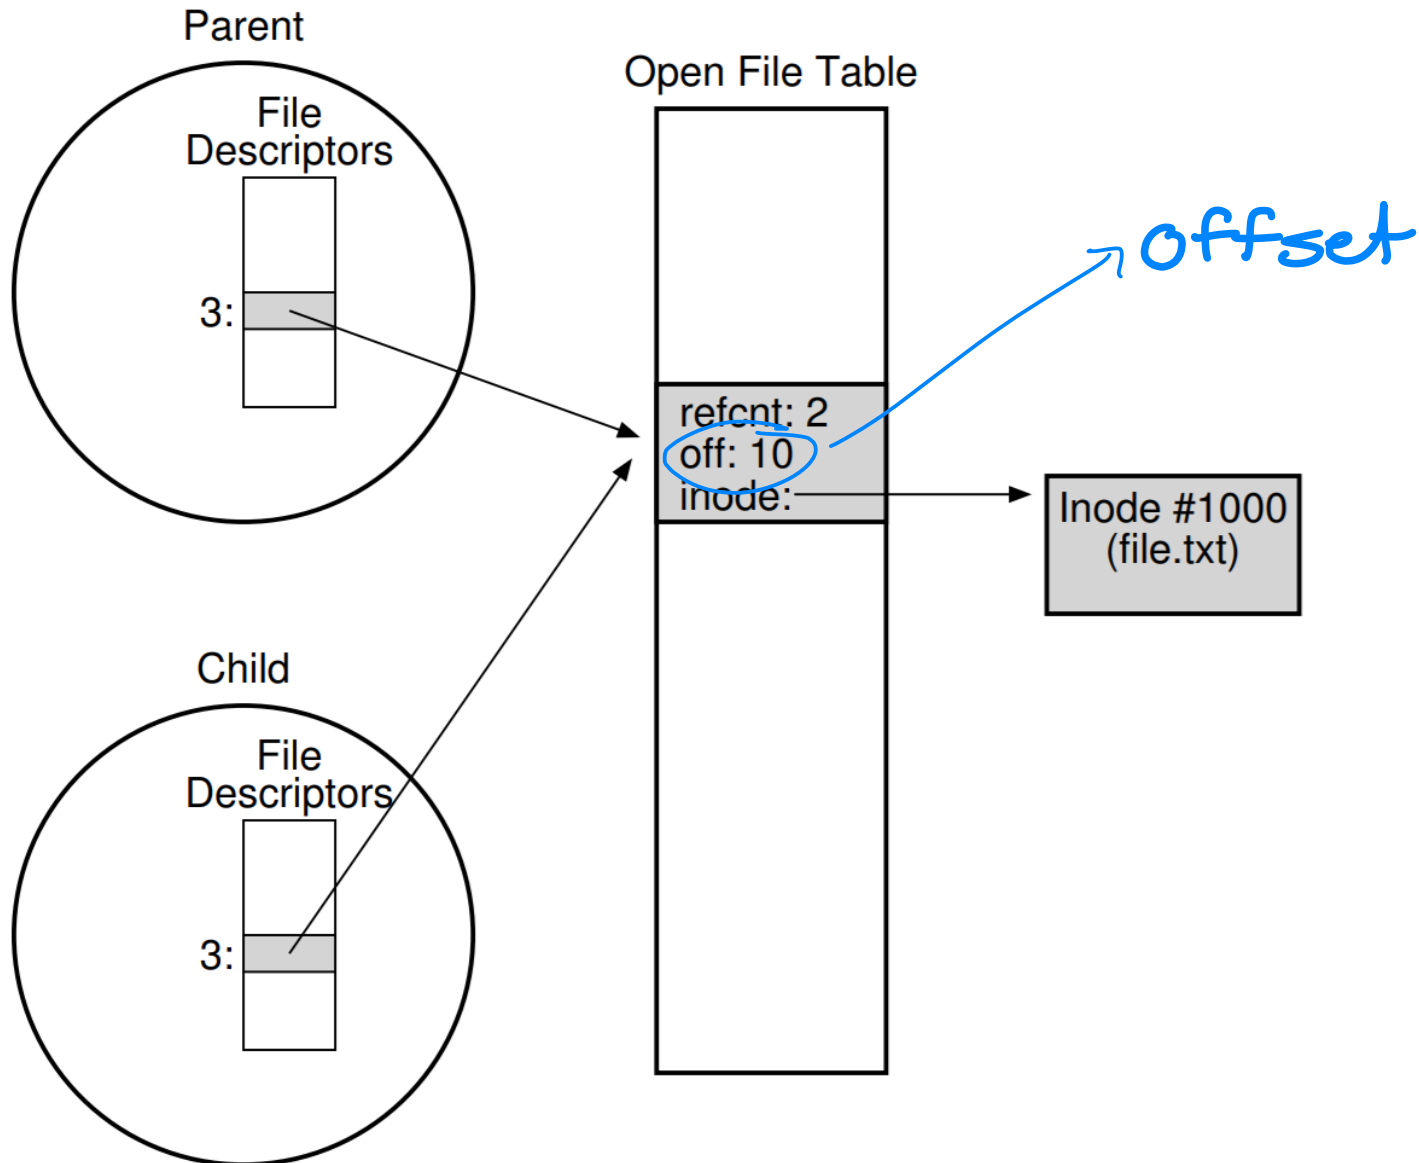
➤ `dup2()`

"`dup()` and friends"

→ file descriptors

inode ~~≠~~ number

File Descriptor => Inode



File System Implementations

In general:

- Pure software
- lots of flexibility
- lots of different solutions

⇒ device characteristics affect
performance greatly

Part 1: Disk

Setup

Disk is composed of **sectors**

- 512-byte block, label $0:n-1$ (n sectors)
- disk is not Byte addressable - **sector addr.**
- Sector writes are atomic

The “unwritten contract”: **accessing 2 sectors near ea is faster**
is faster than 2 sectors far apart ★ **Disk is terrible @ Random access**

Geometry of Disk

Platter: a hard circular surface on which data is stored

Surface: side of a platter (2: front/back)

Track: a concentric circle of sectors, on a surface $\approx 300k$ tracks per surface

Spindle: bounds platters together, and is connected to a motor which spins @ a const rate (rpm)

Disk Head: a mechanism to r/w the magnetic patterns on a track
↳ data

Visual




Figure 37.2: A Single Track Plus A Head


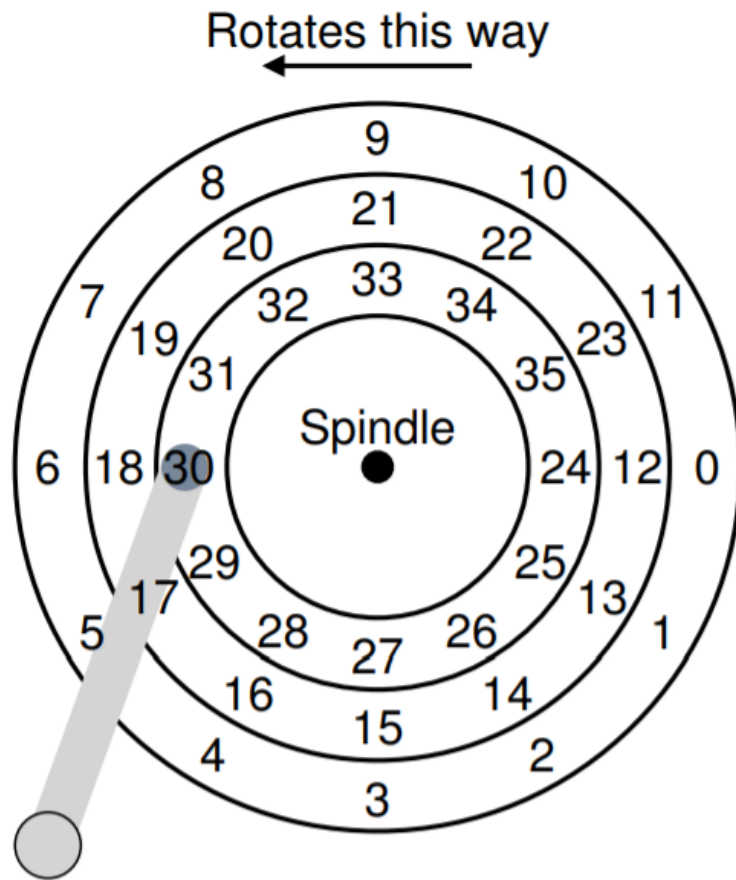



Figure 37.3: Three Tracks Plus A Head (Right: With Seek)

1. "seek": move head to correct track $\sim ms$
2. "rotational delay": wait for disk to spin to the correct position
3. "transfer"

Comparison to RAM



I/O Time

$$T_{I/O} = T_{seek} + T_{rotation} + T_{transfer}$$

how to reduce?

T_{seek} : move tracks as little as possible

T_{rot} : spin faster (not feasible)
anticipate rotation

★
→ embrace
locality

T_{trans} : improve physics